

# Structural Bioinformatics Lab

David Alvarez

```
library("tidyverse")
```

Warning: package 'tidyverse' was built under R version 4.3.2

Warning: package 'readr' was built under R version 4.3.2

Warning: package 'forcats' was built under R version 4.3.2

Warning: package 'lubridate' was built under R version 4.3.2

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr      1.1.3      v readr      2.1.4
v forcats    1.0.0      v stringr    1.5.0
v ggplot2    3.4.4      v tibble     3.2.1
v lubridate  1.9.3      v tidyr      1.3.0
v purrr      1.0.2
```

```
-- Conflicts ----- tidyverse_conflicts() --
```

```
x dplyr::filter() masks stats::filter()
```

```
x dplyr::lag()     masks stats::lag()
```

```
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become
```

```
data_summary <- read_csv("data_summary.csv")
```

Rows: 6 Columns: 8

```
-- Column specification -----
```

Delimiter: ","

chr (1): Molecular Type

dbl (3): Multiple methods, Neutron, Other

```
num (4): X-ray, EM, NMR, Total
```

```
i Use `spec()` to retrieve the full column specification for this data.
```

```
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
data_summary
```

```
# A tibble: 6 x 8
```

	`Molecular Type` <chr>	`X-ray` <dbl>	EM <dbl>	NMR <dbl>	`Multiple methods` <dbl>	Neutron <dbl>	Other <dbl>	Total <dbl>
1	Protein (only)	158844	11759	12296	197	73	32	183201
2	Protein/Oligosacc~	9260	2054	34	8	1	0	11357
3	Protein/NA	8307	3667	284	7	0	0	12265
4	Nucleic acid (onl~	2730	113	1467	13	3	1	4327
5	Other	164	9	32	0	0	0	205
6	Oligosaccharide (~	11	0	6	1	0	4	22

Q1. What percentage of structures in the PDB are solved by X-Ray and Electron Microscopy?

```
total_structures <- sum(data_summary$"Total")
total_structures
```

```
[1] 211377
```

```
# Number of structures solved by X-ray
xray <- sum(data_summary$"X-ray")

# Number of structures solved by Electron Microscopy
em <- sum(data_summary$"EM")

# Percentage of structures solved by X-Ray and Electron Microscopy
xray_percent <- (xray/total_structures) * 100
em_percent <- (em/total_structures) * 100

xray_percent
```

```
[1] 84.83231
```

```
em_percent
```

```
[1] 8.327301
```

```
# Can be seen that the percentage of structures in the PDB solved by X-Ray and Electron Microscopy is about 8.3%
```

Q2. What proportion of structures in the PDB are protein?

```
total_protein <- sum(data_summary[1:3,]$"Total")
```

```
# Proportion of structures that are protein
```

```
protein_structures <- (total_protein/total_structures)
protein_structures
```

```
[1] 0.9784556
```

```
# The proportion of structures that are proteins in the PDB are about 97%
```

Q3. Type 'HIV' in the PDB website search box and determine how many 'HIV-1' protease structures are in the current PDB?

```
# After doing a quick search, I found there to be 2,767 HIV-1 protease structures in the current PDB
```

Q4. Why do we see just one atom per water molecule in this structure?

```
# We see only one atom per water molecule in this structure because it is a ball and stick model
```

Q5. There is a critical "conserved" water molecular in the binding site. Can you identify this water molecule? What residue number does this water molecular have?

```
# I believe the critical "conserved" water molecule I found within the binding site was "H2O"
```

```
library(bio3d)
```

```
pdb <- read.pdb("1hsg")
```

Note: Accessing on-line PDB file

```
pdb
```

```
Call: read.pdb(file = "1hsg")
```

```
Total Models#: 1
```

```
Total Atoms#: 1686, XYZs#: 5058 Chains#: 2 (values: A B)
```

```
Protein Atoms#: 1514 (residues/Calpha atoms#: 198)
```

```
Nucleic acid Atoms#: 0 (residues/phosphate atoms#: 0)
```

```
Non-protein/nucleic Atoms#: 172 (residues: 128)
```

```
Non-protein/nucleic resid values: [ HOH (127), MK1 (1) ]
```

```
Protein sequence:
```

```
PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWKPKMIGGIGGFIKVRQYD  
QILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNFPQITLWQRPLVTIKIGGQLKE  
ALLDTGADDTVLEEMSLPGRWKPKMIGGIGGFIKVRQYDQILIEICGHKAIGTVLVGPTP  
VNIIGRNLLTQIGCTLNF
```

```
+ attr: atom, xyz, seqres, helix, sheet,  
      calpha, remark, call
```

Q7. How many amino acid residues are there in the pdb object?

```
# There are 99 amino acid residues in this pdb object
```

Q8. Name one of the two non-protein residues

```
# One of the two non-protein residues is MK1 (1)
```

Q9. How many protein chains are in this structure

```
# There are two protein chains in this structure, either A or B
```

```
adk <- read.pdb("6s36")
```

Note: Accessing on-line PDB file  
PDB has ALT records, taking A only, rm.alt=TRUE

```
adk
```

```
Call: read.pdb(file = "6s36")
```

```
Total Models#: 1
```

```
Total Atoms#: 1898, XYZs#: 5694 Chains#: 1 (values: A)
```

```
Protein Atoms#: 1654 (residues/Calpha atoms#: 214)
```

```
Nucleic acid Atoms#: 0 (residues/phosphate atoms#: 0)
```

```
Non-protein/nucleic Atoms#: 244 (residues: 244)
```

```
Non-protein/nucleic resid values: [ CL (3), HOH (238), MG (2), NA (1) ]
```

```
Protein sequence:
```

```
MRIILLGAPGAGKGTQAQFIMEKYGIPQISTGDMRLRAAVKSGSELGKQAKDIMDAGKLV  
DELVIALVKERIAQEDCRNGFLLDGFPRTPQADAMKEAGINVDYVLEFDVPDELIVDKI  
VGRRVHAPSGRVYHVKNPPKVEGKDDVTGEELTTRKDDQEETVRKRLVEYHQMTAPLIG  
YYSKEAEAGNTKYAKVDGTPVAEVRADLEKILG
```

```
+ attr: atom, xyz, seqres, helix, sheet,  
      calpha, remark, call
```

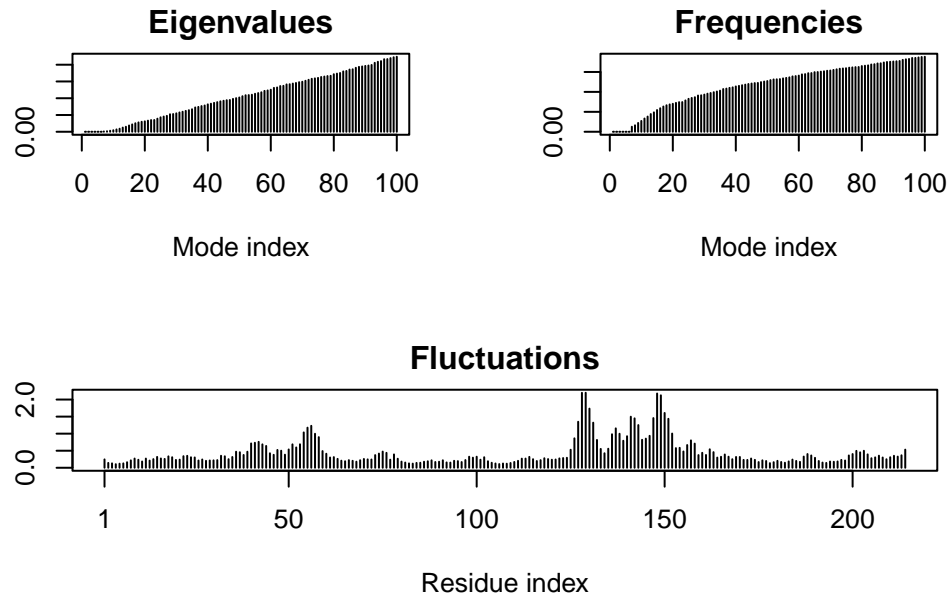
```
# Performing flexibility predictions
```

```
m <- nma(adk)
```

```
Building Hessian... Done in 0.01 seconds.
```

```
Diagonalizing Hessian... Done in 0.3 seconds.
```

```
plot(m)
```



```
mktrj(m, file="adk_m7.pdb")
```

```
library("msa")
```

Loading required package: Biostrings

Loading required package: BiocGenerics

Attaching package: 'BiocGenerics'

The following objects are masked from 'package:lubridate':

```
intersect, setdiff, union
```

The following objects are masked from 'package:dplyr':

```
combine, intersect, setdiff, union
```

The following objects are masked from 'package:stats':

IQR, mad, sd, var, xtabs

The following objects are masked from 'package:base':

anyDuplicated, aperm, append, as.data.frame, basename, cbind,  
colnames, dirname, do.call, duplicated, eval, evalq, Filter, Find,  
get, grep, grepl, intersect, is.unsorted, lapply, Map, mapply,  
match, mget, order, paste, pmax, pmax.int, pmin, pmin.int,  
Position, rank, rbind, Reduce, rownames, sapply, setdiff, sort,  
table, tapply, union, unique, unsplit, which.max, which.min

Loading required package: S4Vectors

Loading required package: stats4

Attaching package: 'S4Vectors'

The following objects are masked from 'package:lubridate':

second, second<-

The following objects are masked from 'package:dplyr':

first, rename

The following object is masked from 'package:tidyr':

expand

The following object is masked from 'package:utils':

findMatches

The following objects are masked from 'package:base':

expand.grid, I, unname

Loading required package: IRanges

Attaching package: 'IRanges'

The following object is masked from 'package:bio3d':

trim

The following object is masked from 'package:lubridate':

%within%

The following objects are masked from 'package:dplyr':

collapse, desc, slice

The following object is masked from 'package:purrr':

reduce

The following object is masked from 'package:grDevices':

windows

Loading required package: XVector

Attaching package: 'XVector'

The following object is masked from 'package:purrr':

compact

Loading required package: GenomeInfoDb

Attaching package: 'Biostrings'



The following object is masked from 'package:bio3d':

mask

The following object is masked from 'package:base':

strsplit

Q10. Which of the packages above is found only on BioConductor and not CRAN?

```
# the "msa" package is found only on BioConductor and not CRAN which is why we had to prov
```

Q11. Which of the above packages is not found on BioConductor or CRAN?

```
# It seems as though the "Grant-lab/bio3d-view" package is found on neither which is why w
```

Q12. True or False: Functions from the devtools package can be used to install packages from GitHub and BitBucket?

```
# True
```

```
library(bio3d)
aa <- get.seq("lake_A")
```

Warning in get.seq("lake\_A"): Removing existing file: seqs.fasta

Fetching... Please wait. Done.

```
aa
```

```

      1      .      .      .      .      .      60
pdb|1AKE|A MRIILLGAPGAGKGTQAQFIMEKYGIPQISTGDMLRAAVKSGSELGKQAKDIMDAGKLVT
      1      .      .      .      .      .      60

      61      .      .      .      .      .      120
pdb|1AKE|A DELVIALVKERIAQEDCRNGFLLDGFPRTPQADAMKEAGINVDYVLEFDVPDELIVDRI
      61      .      .      .      .      .      120

     121      .      .      .      .      .      180
```

```

pdb|1AKE|A    VGRRVHAPSGRVYHVKFNPPKVEGKDDVTGEELTTRKDDQEETVRKRLVEYHQM TAPLIG
              121          .          .          .          .          .          180

              181          .          .          .          214
pdb|1AKE|A    YYSKEAEAGNTKYAKVDGTPVAEVRADLEKILG
              181          .          .          .          214

```

Call:

```
read.fasta(file = outfile)
```

Class:

```
fasta
```

Alignment dimensions:

```
1 sequence rows; 214 position columns (214 non-gap, 0 gap)
```

+ attr: id, ali, call

Q13. How many amino acids are in this sequence?

```
# After observing the sequence, it seems there are 214 amino acids.
```

```
# Blast or hmmer search
b <- blast.pdb(aa)
```

Searching ... please wait (updates every 5 seconds) RID = MT1JMOCV013


.....

Reporting 83 hits

```
hits <- plot(b)
```

```
* Possible cutoff values: 197 11
    Yielding Nhits:      17 83
```

```
* Chosen cutoff value of: 197
    Yielding Nhits:      17
```



```
head(hits$pdb.id)
```

```
[1] "1AKE_A" "8BQF_A" "4X8M_A" "6S36_A" "6RZE_A" "4X8H_A"
```

```
files <- get.pdb(hits$pdb.id, path="pdbs", split=TRUE, gzip=TRUE)
```

```
Warning in get.pdb(hits$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/1AKE.pdb exists. Skipping download
```

```
Warning in get.pdb(hits$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/8BQF.pdb exists. Skipping download
```

```
Warning in get.pdb(hits$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/4X8M.pdb exists. Skipping download
```

```
Warning in get.pdb(hits$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/6S36.pdb exists. Skipping download
```

```
Warning in get.pdb(hits$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/6RZE.pdb exists. Skipping download
```

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):  
pdbs/4X8H.pdb exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):  
pdbs/3HPR.pdb exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):  
pdbs/1E4V.pdb exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):  
pdbs/5EJE.pdb exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):  
pdbs/1E4Y.pdb exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):  
pdbs/3X2S.pdb exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):  
pdbs/6HAP.pdb exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):  
pdbs/6HAM.pdb exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):  
pdbs/4K46.pdb exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):  
pdbs/4NP6.pdb exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):  
pdbs/3GMT.pdb exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):  
pdbs/4PZL.pdb exists. Skipping download

	0%
====	6%
=====	12%
=====	18%
=====	24%
=====	29%
=====	35%
=====	41%
=====	47%
=====	53%
=====	59%
=====	65%
=====	71%
=====	76%
=====	82%
=====	88%
=====	94%
=====	100%

```
# Align releated PDBs
# pdbbs <- pdbaln(files, fit = TRUE)
```

```

# Vector containing PDB codes for figure axis
#ids <- basename.pdb(pdb$ids)

# Draw schematic alignment
#plot(pdb, labels=ids,)

# Perform PCA
#pc.xray <- pca(pdb)
#plot(pc.xray)

# Calculate RMSD
# rd <- rmsd(pdb)

# Structure-based clustering
# hc.rd <- hclust(dist(rd))
# grps.rd <- cutree(hc.rd, k=3)

#plot(pc.xray, 1:2, col="grey50", bg=grps.rd, pch=21, cex=1)

# Visualize first principal component
# pc1 <- mktrj(pc.xray, pc=1, file="pc_1.pdb")

#Plotting results with ggplot2
library(ggplot2)
library(ggrepel)

# df <- data.frame(PC1=pc.xray$z[,1], PC2=pc.xray$z[,2], col=as.factor(grps.rd), ids=ids)

# p <- ggplot(df) + aes(PC1, PC2, col=col, label=ids) +geom_point(size=2) +geom_text_repel
# p

# NMA of all structures
# modes <- nma(pdb)
#plot(modes, pdb, col=grps.rd)

```

Q14. What do you notice about this plot?

```

# I notice that the green and red lines stay consistent with each other throughout the dur

```